Robust Retrieval-Augmented Generation with Source Citation and Hallucination Mitigation

Patrick Lewis, Maria Rodriguez, Kevin Patel
Center for AI Safety and Information Retrieval, 2024

ABSTRACT
Large language models (LLMs) deployed for scientific inquiry frequently hallucinate plausible-sounding but completely fabricated assertions, citations, and experimental outcomes. Retrieval-Augmented Generation (RAG) mitigates this vulnerability by conditioning text generation on retrieved external documents. However, conventional RAG systems suffer from citation misalignment, where generated claims attribute findings to irrelevant or non-existent document pages. In this paper, we present VeriRAG, a robust retrieval framework incorporating hybrid dense-sparse retrieval via Reciprocal Rank Fusion (RRF) and post-generation citation mapping. Across rigorous evaluation benchmarks, VeriRAG reduces ungrounded hallucinations by 82.5% and achieves a citation attribution precision of 94.1%.

1. INTRODUCTION
In high-stakes academic and clinical domains, ungrounded AI answers undermine user trust and propagate misinformation. When asked about experimental findings or methodological constraints, standard parametric models generate confident statements with fictitious authors and non-existent DOIs. RAG architectures offer an appealing solution by decoupling knowledge storage from parametric weights. Yet, naive chunking and basic vector similarity search frequently retrieve irrelevant snippets, leading the generator to hallucinate bridging explanations between mismatched fragments.

2. HYBRID RETRIEVAL ARCHITECTURE
VeriRAG employs a dual-stage retrieval pipeline combining dense vector embeddings with lexical BM25 retrieval. While dense bi-encoders excel at capturing thematic semantic similarity, they frequently miss exact technical terms, mathematical variable symbols, or specific author surnames. Conversely, BM25 enforces exact keyword precision but fails on paraphrased natural language questions.
We unify dense and sparse rankings using Reciprocal Rank Fusion (RRF):
RRF_Score(d) = (w_dense / (k + rank_dense(d))) + (w_sparse / (k + rank_sparse(d)))
where k=60, w_dense=0.65, and w_sparse=0.35.

3. CITATION MAPPING AND VERIFICATION ENGINE

Unlike standard chatbots that append generic bibliographies, VeriRAG implements granular provenance tracking. Every retrieved passage is tagged with an immutable source identifier containing document ID, page number, and section name.

During generation, the model is constrained by system prompts to tag each factual claim with an exact bracketed index [X].

Our citation verification engine performs two critical checks:

1. Lexical and NLI entailment: Does the cited chunk contain semantic entailment for the attributed assertion?
2. Page Provenance Validation: Verifies that the cited page exists in the physical PDF and maps the exact coordinates for interactive document viewer highlighting.

If the similarity score of retrieved passages falls below threshold tau=0.20, or if evidence is insufficient, the system enforces a strict refusal protocol: "I could not find sufficient evidence in the uploaded papers to answer this question."

4. EXPERIMENTAL EVALUATION

We evaluated VeriRAG against Dense-Only RAG, Sparse-Only RAG, and Closed-Book GPT-4 across 1,200 scientific questions drawn from PubMed QA and the SciCite benchmark.

Evaluation metrics comprise:

- Retrieval Precision@5
- Faithfulness (percentage of claims verified by cited source passages)
- Citation Coverage (percentage of non-trivial claims supported by citations)
- Latency (milliseconds per answer)

5. RESULTS AND ABLATION STUDY

Experimental results reveal that VeriRAG achieves 94.1% faithfulness, compared to only 68.2% for Dense-Only RAG and 31.4% for Closed-Book LLMs.

The combination of RRF hybrid retrieval improved Retrieval Precision@5 from 0.71 to 0.89.

Ablation studies demonstrate that removing BM25 lexical matching caused a 28% drop in retrieval accuracy for queries containing specialized chemical formulas or algorithm acronyms.

6. LIMITATIONS

We note two primary limitations:

First, multi-hop reasoning requiring synthesis across more than four separate papers occasionally struggles with context window fragmentation.

Second, the latency overhead of hybrid retrieval and reranking adds approximately 120ms compared to unverified greedy generation, although this trade-off is highly justified for academic integrity.

7. CONCLUSION

VeriRAG demonstrates that disciplined hybrid retrieval, strict citation attribution, and automated refusal safeguards virtually eliminate hallucinated claims in scientific question answering.

REFERENCES

[1] Lewis et al., Retrieval-augmented generation for knowledge-intensive NLP tasks. NeurIPS 2020.

[2] Robertson et al., The probabilistic relevance framework: BM25 and beyond. FnTIR 2009.

[3] Gao et al., Retrieval-augmented generation for large language models: A survey. arXiv:2312.10997, 2023.